


Efficiency on Qwen3-VL-4B (single 24G GPU, bf16)

Time-to-first-token vs. compression



Decoding throughput vs. compression

